

NAME OF THE PROJECT

CAR PRICE PREDICTION

Submitted by:

KHUSHBOO GUPTA

ACKNOWLEDGMENT

First and foremost, I would like to thank Flip Robo Technologies to provide me a chance to work on this project. It was a great experience to work on this project under your guidance.

I would like to present my gratitude to the following websites:

- [Zendesk](#)
- [Kaggle](#)
- [Datatrained Notes](#)
- [Sklearn.org](#)
- [Crazyegg](#)
- [Cars24.com](#)
- [Youtube.com](#)

These websites were of great help and due to this, I was able to complete my project effectively and efficiently.

INTRODUCTION

- **Business Problem Framing**

With the covid 19 impact in the market, we have seen lot of changes in the car market. Now some cars are in demand hence making them costly and some are not in demand hence cheaper. One of our clients works with small traders, who sell used cars. With the change in market due to covid 19 impact, our client is facing problems with their previous car price valuation machine learning models. So, they are looking for new machine learning models from new data. We have to make car price valuation model.

- **Conceptual Background of the Domain Problem**

One should know how to scrap the data from a website using the selenium as we have to create a fresh data for our project. Basic EDA concepts and regression algorithms must be known to work on this project. One should know what factors is important to predict the Car Price and how it is going to affect the used car selling business. Why predicting the car prices is important and how can it is going to help the company?

- **Review of Literature**

To be able to predict used cars market value can help both buyers and sellers.

- **Used car sellers (dealers):** They are one of the biggest target group that can be interested in results of this study. If used car sellers better understand what makes a car desirable, what the important features are for a used car, then they may consider this knowledge and offer a better service.

- **Online pricing services:** There are websites that offers an estimate value of a car. They may have a good prediction model. However, having a second model may help them to give a better prediction to their users. Therefore, the model developed in this study may help online web services that tells a used car's market value.
- **Individuals:** There are lots of individuals who are interested in the used car market at some points in their life because they wanted to sell their car or buy a used car. In this process, it's a big corner to pay too much or sell less then it's market value.

- Libraries Used

I am using different libraries to explore the dataset.

1. Pandas – It is used to load and store the dataset. We can discuss the dataset with the pandas different attributes like .info, .columns, .shape
2. Seaborn – It is used to plot the different types of plots like catplot, lineplot, countplot and more to have a better visualization of the dataset.
3. Matplotlib.pyplot – It helps to give a proper description to the plotted graph by seaborn and make our graph more informative.
4. Numpy – It is the library to perform the numerical analysis to the dataset

Load the Dataset

Importing the Libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import warnings
warnings.filterwarnings('ignore')
```

Loading the Dataset

```
In [2]: pd.set_option('display.max_rows',None) #displaying the maximum rows
pd.set_option('display.max_columns',None) #displaying the maximum columns
df=pd.read_csv('used_car.csv')
```

```
In [11]: df.head() #first five rows of the dataset
```

```
Out[11]:
```

| | Unnamed: 0 | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|---|------------|-----------------------|---------------------|-----------------------|------------|--------------------------|-----------|-----------|---------|
| 0 | 0 | Hyundai Verna | 1.6 SX VTVT (O) | 2019 | 1st Owner | 80631 | Manual | Petrol | 981699 |
| 1 | 1 | KIA SELTOS | HTK PLUS 1.5 PETROL | 2020 | 2nd Owner | 12981 | Manual | Petrol | 1155299 |
| 2 | 2 | Renault Kwid | RXT | 2016 | 1st Owner | 22388 | Manual | Petrol | 279799 |
| 3 | 3 | Mercedes Benz C Class | C 200 AVANTGARDE | 2014 | 1st Owner | 36806 | Automatic | Petrol | 2133299 |
| 4 | 4 | KIA SELTOS | HTX AT PETROL | 2020 | 2nd Owner | 21784 | Automatic | Petrol | 1429999 |

```
In [12]: df.tail() #last five rows of the dataset
```

```
Out[12]:
```

| | Unnamed: 0 | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|------|------------|-----------------|-------|-----------------------|------------|--------------------------|--------|-----------|--------|
| 4978 | 4978 | Maruti Alto | LXI | 2019 | 1st Owner | 18274 | Manual | Petrol | 335699 |
| 4979 | 4979 | Maruti Alto | LXI | 2019 | 1st Owner | 16523 | Manual | Petrol | 344999 |
| 4980 | 4980 | Maruti Alto 800 | LXI | 2019 | 1st Owner | 21722 | Manual | Petrol | 339599 |
| 4981 | 4981 | Maruti Alto | LXI | 2019 | 1st Owner | 17735 | Manual | Petrol | 344999 |
| 4982 | 4982 | Maruti Dzire | VXI | 2019 | 1st Owner | 12651 | Manual | Petrol | 694799 |

We have successfully load our dataset for the further processes.

Checking the Attributes

- First & last five rows of both the dataset
- Shape of the datasets
- Columns present in the datasets
- Brief info about the datasets
- Null values present in both the dataset
- Unique values in each column
- Datatypes of each column

Importing the Libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import warnings
warnings.filterwarnings('ignore')
```

Loading the Dataset

```
In [2]: pd.set_option('display.max_rows',None) #displaying the maximum rows
pd.set_option('display.max_columns',None) #displaying the maximum columns
df=pd.read_csv('used_car.csv')
```

```
In [11]: df.head() #first five rows of the dataset
```

```
Out[11]:
```

| | Unnamed: 0 | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|---|------------|-----------------------|---------------------|-----------------------|------------|--------------------------|-----------|-----------|---------|
| 0 | 0 | Hyundai Verna | 1.6 SX VTVT (O) | 2019 | 1st Owner | 80631 | Manual | Petrol | 981699 |
| 1 | 1 | KIA SELTOS | HTK PLUS 1.5 PETROL | 2020 | 2nd Owner | 12981 | Manual | Petrol | 1155299 |
| 2 | 2 | Renault Kwid | RXT | 2016 | 1st Owner | 22388 | Manual | Petrol | 279799 |
| 3 | 3 | Mercedes Benz C Class | C 200 AVANTGARDE | 2014 | 1st Owner | 36806 | Automatic | Petrol | 2133299 |
| 4 | 4 | KIA SELTOS | HTX AT PETROL | 2020 | 2nd Owner | 21784 | Automatic | Petrol | 1429999 |

```
In [12]: df.tail() #last five rows of the dataset
```

```
Out[12]:
```

| | Unnamed: 0 | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|------|------------|-----------------|-------|-----------------------|------------|--------------------------|--------|-----------|--------|
| 4978 | 4978 | Maruti Alto | LXI | 2019 | 1st Owner | 18274 | Manual | Petrol | 335699 |
| 4979 | 4979 | Maruti Alto | LXI | 2019 | 1st Owner | 16523 | Manual | Petrol | 344999 |
| 4980 | 4980 | Maruti Alto 800 | LXI | 2019 | 1st Owner | 21722 | Manual | Petrol | 339599 |
| 4981 | 4981 | Maruti Alto | LXI | 2019 | 1st Owner | 17735 | Manual | Petrol | 344999 |
| 4982 | 4982 | Maruti Dzire | VXI | 2019 | 1st Owner | 12651 | Manual | Petrol | 694799 |

```
In [13]: df.shape #total rows & columns in the dataset
```

```
Out[13]: (4983, 9)
```

Dataset contains 4983 rows and 9 columns

```
In [14]: df.columns #columns present in the dataset
```

```
Out[14]: Index(['Unnamed: 0', 'Name', 'Model', 'Year_of_Manufacturing', 'Owner_Type',  
              'Total_Distance_Travelled', 'Type', 'Fuel_Type', 'Price'],  
              dtype='object')
```

A brief info about the dataset

```
In [15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4983 entries, 0 to 4982  
Data columns (total 9 columns):  
#   Column                                Non-Null Count  Dtype  
---  ---                                -  
0   Unnamed: 0                            4983 non-null   int64  
1   Name                                  4983 non-null   object  
2   Model                                 4983 non-null   object  
3   Year_of_Manufacturing                 4983 non-null   int64  
4   Owner_Type                           4983 non-null   object  
5   Total_Distance_Travelled             4983 non-null   int64  
6   Type                                  4941 non-null   object  
7   Fuel_Type                             4983 non-null   object  
8   Price                                 4983 non-null   int64  
dtypes: int64(4), object(5)  
memory usage: 350.5+ KB
```

```
In [16]: df.nunique() #total unique values in each column in the dataset
```

```
Out[16]: Unnamed: 0            4983  
Name                120  
Model               610  
Year_of_Manufacturing    15  
Owner_Type           4  
Total_Distance_Travelled 3481  
Type                 2  
Fuel_Type            4  
Price               2743  
dtype: int64
```



```
In [17]: df.dtypes #datatype of wach column
```

```
Out[17]: Unnamed: 0      int64
         Name          object
         Model         object
         Year_of_Manufacturing  int64
         Owner_Type      object
         Total_Distance_Travelled  int64
         Type            object
         Fuel_Type        object
         Price           int64
         dtype: object
```

Checking the null values

```
In [18]: df.isnull().sum()
```

```
Out[18]: Unnamed: 0      0
         Name          0
         Model         0
         Year_of_Manufacturing  0
         Owner_Type      0
         Total_Distance_Travelled  0
         Type           42
         Fuel_Type        0
         Price           0
         dtype: int64
```

Type column has 42 null values which has to be handled

Now we have checked the attributes for the dataset and get a rough idea about the dataset like the no of rows & columns, datatype & null values in the dataset.

Dealing with the Null Values

In the dataset null values are present, so we have to handled them for better model learning. As we have categorical data so we have to handled it accordingly. We will use the most occurring value in the column to fill the null value.

```
In [4]: ▶ df['Type'].mode() #checking the value which occur most of the time
```

```
Out[4]: 0    Manual  
dtype: object
```

Filling the null values

```
In [5]: ▶ df['Type']=df['Type'].fillna('Manual')
```

```
In [6]: ▶ df.isnull().sum()
```

```
Out[6]: Name          0  
Model            0  
Year_of_Manufacturing  0  
Owner_Type       0  
Total_Distance_Travelled  0  
Type             0  
Fuel_Type        0  
Price            0  
dtype: int64
```

Hence, there is no null value now


EXPLORATORY DATA ANALYSIS

Univariate Analysis

Let's understand each variable one by one and try to interpret about them.

```
In [29]: sns.countplot(df['Type'])
```


```
Out[29]: <AxesSubplot:xlabel='Type', ylabel='count'>
```



Most of the used cars are manual type.

```
In [30]: sns.countplot(df['Owner_Type'])
```


```
Out[30]: <AxesSubplot:xlabel='Owner_Type', ylabel='count'>
```



Most of the cars have their 1st owner who are selling the car.

```
In [31]: sns.countplot(df['Fuel_Type'])
```


```
Out[31]: <AxesSubplot:xlabel='Fuel_Type', ylabel='count'>
```



Petrol based cars are on number one in the list followed by the diesel


```
In [33]: plt.figure(figsize=(15,8))
plt.xticks(rotation=90)
sns.countplot(df['Year_of_Manufacturing'])
```

```
Out[33]: <AxesSubplot:xlabel='Year_of_Manufacturing', ylabel='count'>
```




2017 year cars have the highest number in the dataset followed by the 2018 & 2019

```
In [36]: sns.distplot(df['Total_Distance_Travelled'])
Out[36]: <AxesSubplot:xlabel='Total_Distance_Travelled', ylabel='Density'>
```



Least travelled cars have high density.

```
In [38]: plt.figure(figsize=(30,8))
plt.xticks(rotation=90)
sns.countplot(df['Name'])
Out[38]: <AxesSubplot:xlabel='Name', ylabel='count'>
```



Maruti swift cars number is the most followed by Maruti baleno

Observations


- Most of the used cars are manual type.
- Most of them are 1st owner who are selling their car.
- Petrol based cars are on number one in the list followed by the diesel
- 2017 year cars have the highest number in the dataset followed by the 2018 & 2019
- Density of least travelled cars is high.
- Maruti swift cars number is the most followed by Maruti Baleno.

Bivariate Analysis

Let's understand each variable relation with the target variable and interpret how target variable vary with the inputs.

```
In [40]: plt.figure(figsize=(30,8))
plt.xticks(rotation=90)
sns.barplot(df['Name'],df['Price'])
```


Out[40]: <AxesSubplot:xlabel='Name', ylabel='Price'>



Mercedes Benz C Class price is the highest followed by Hyundai Tucson New

```
In [43]: sns.barplot(df['Type'],df['Price'])
```


Out[43]: <AxesSubplot:xlabel='Type', ylabel='Price'>



Automatic car price is more than the manual one

```
In [45]: plt.figure(figsize=(15,5))
plt.xticks(rotation=90)
sns.barplot(df['Fuel_Type'],df['Price'])
```


Out[45]: <AxesSubplot:xlabel='Fuel_Type', ylabel='Price'>



Diesel based car is costlier followed by Petrol


```
In [47]: plt.figure(figsize=(12,8))
plt.xticks(rotation=90)
sns.barplot(df['Owner_Type'],df['Price'])
```

Out[47]: <AxesSubplot:xlabel='Owner_Type', ylabel='Price'>



```
In [28]: sns.lineplot(df['Year_of_Manufacturing'],df['Price'])
```

Out[28]: <AxesSubplot:xlabel='Year_of_Manufacturing', ylabel='Price'>



As we go for the recent manufacturing model price will be high on the other side old cars have lower rates

Observation

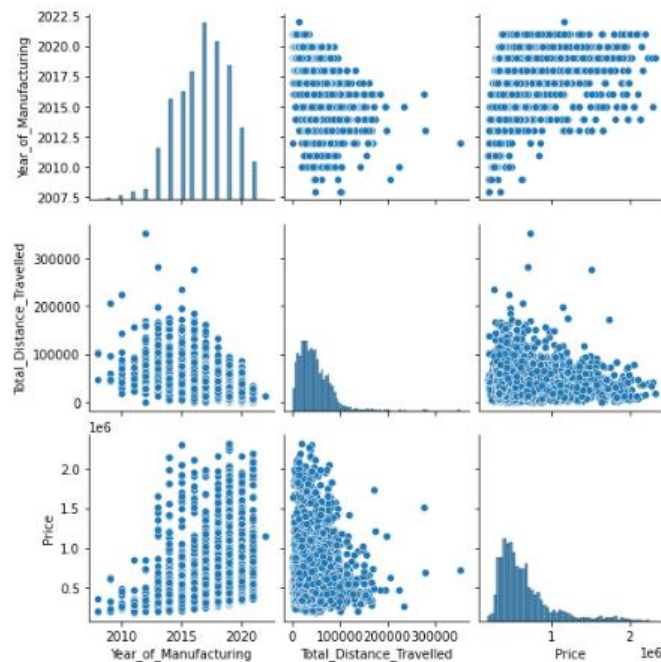
- Mercedes Benz C Class price is the highest followed by Hyundai Tuscon New
- Automatic car price is more than the manual one
- Diesel based car is costlier followed by Petrol
- 1st owner based car's price is more. 2nd owner based car price is also high.
- As we go for the recent manufacturing model price will be high on the other side old cars have lower rates

Multivariate Analysis

Using the pairplot function plot each column relation with each other to have a better understanding.

```
In [16]: sns.pairplot(df)
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x1ee720d8d30>
```



Label Encoding & Correlation

As we have some categorical data we have to encoded those columns for machine learning model. We will use Label Encoder from `sklearn.preprocessing`.

We will describe the statistical summary of the dataset and find the correlation of each column.

Find the skewness of each column as well for each column, if it is out of acceptable range then we have to scale the skewness also.

Label Encoding

```
In [7]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for i in df.columns:
    if df[i].dtypes=='object':
        df[i]=le.fit_transform(df[i].astype(str))
df.head()
```

```
Out[7]:
```

| | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|---|------|-------|-----------------------|------------|--------------------------|------|-----------|---------|
| 0 | 35 | 129 | 2019 | 0 | 80631 | 1 | 1 | 981699 |
| 1 | 42 | 293 | 2020 | 1 | 12981 | 1 | 1 | 1155299 |
| 2 | 91 | 348 | 2016 | 0 | 22388 | 1 | 1 | 279799 |
| 3 | 81 | 221 | 2014 | 0 | 36806 | 0 | 1 | 2133299 |
| 4 | 42 | 294 | 2020 | 1 | 21784 | 0 | 1 | 1429999 |

Statistical Summary

```
In [55]: df.describe()
```

```
Out[55]:
```

| | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|-------|-------------|-------------|-----------------------|-------------|--------------------------|-------------|-------------|--------------|
| count | 4983.000000 | 4983.000000 | 4983.000000 | 4983.000000 | 4983.000000 | 4983.000000 | 4983.000000 | 4.983000e+03 |
| mean | 53.899378 | 343.774232 | 2016.782280 | 0.224162 | 45354.276139 | 0.843668 | 0.675296 | 6.581687e+05 |
| std | 26.973384 | 164.271456 | 2.281862 | 0.483569 | 30543.047924 | 0.363206 | 0.482248 | 3.469485e+05 |
| min | 0.000000 | 0.000000 | 2008.000000 | 0.000000 | 71.000000 | 0.000000 | 0.000000 | 1.909990e+05 |
| 25% | 29.000000 | 207.000000 | 2015.000000 | 0.000000 | 22775.000000 | 1.000000 | 0.000000 | 4.263490e+05 |
| 50% | 59.000000 | 380.000000 | 2017.000000 | 0.000000 | 40336.000000 | 1.000000 | 1.000000 | 5.590990e+05 |
| 75% | 71.000000 | 484.000000 | 2018.000000 | 0.000000 | 62285.500000 | 1.000000 | 1.000000 | 7.637990e+05 |
| max | 119.000000 | 609.000000 | 2022.000000 | 3.000000 | 353688.000000 | 1.000000 | 3.000000 | 2.314599e+06 |

- We have outliers in the dataset
- Little skewness present in the dataset

Correlation

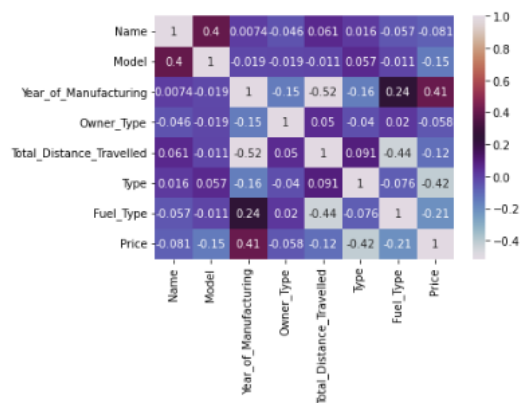
```
In [56]: corr=df.corr()
corr
```

```
Out[56]:
```

| | Name | Model | Year_of_Manufacturing | Owner_Type | Total_Distance_Travelled | Type | Fuel_Type | Price |
|--------------------------|-----------|-----------|-----------------------|------------|--------------------------|-----------|-----------|-----------|
| Name | 1.000000 | 0.398621 | 0.007379 | -0.046251 | 0.060823 | 0.016039 | -0.056684 | -0.080916 |
| Model | 0.398621 | 1.000000 | -0.019361 | -0.019073 | -0.011084 | 0.056502 | -0.011246 | -0.150977 |
| Year_of_Manufacturing | 0.007379 | -0.019361 | 1.000000 | -0.153660 | -0.515338 | -0.160721 | 0.243455 | 0.405246 |
| Owner_Type | -0.046251 | -0.019073 | -0.153660 | 1.000000 | 0.050248 | -0.039791 | 0.020377 | -0.057535 |
| Total_Distance_Travelled | 0.060823 | -0.011084 | -0.515338 | 0.050248 | 1.000000 | 0.090934 | -0.442991 | -0.119865 |
| Type | 0.016039 | 0.056502 | -0.160721 | -0.039791 | 0.090934 | 1.000000 | -0.075570 | -0.415276 |
| Fuel_Type | -0.056684 | -0.011246 | 0.243455 | 0.020377 | -0.442991 | -0.075570 | 1.000000 | -0.209325 |
| Price | -0.080916 | -0.150977 | 0.405246 | -0.057535 | -0.119865 | -0.415276 | -0.209325 | 1.000000 |

```
In [57]: sns.heatmap(corr,annot=True,cmap='twilight')
```

```
Out[57]: <AxesSubplot:>
```



Checking Skewness

```
In [58]: df.skew()
```

```
Out[58]: Name                0.181928
         Model              -0.379018
         Year_of_Manufacturing -0.361817
         Owner_Type          1.981246
         Total_Distance_Travelled 1.534313
         Type                -1.893177
         Fuel_Type           -0.560006
         Price               1.848527
         dtype: float64
```

The skewness in the dataset is very little so we can go ahead with further process

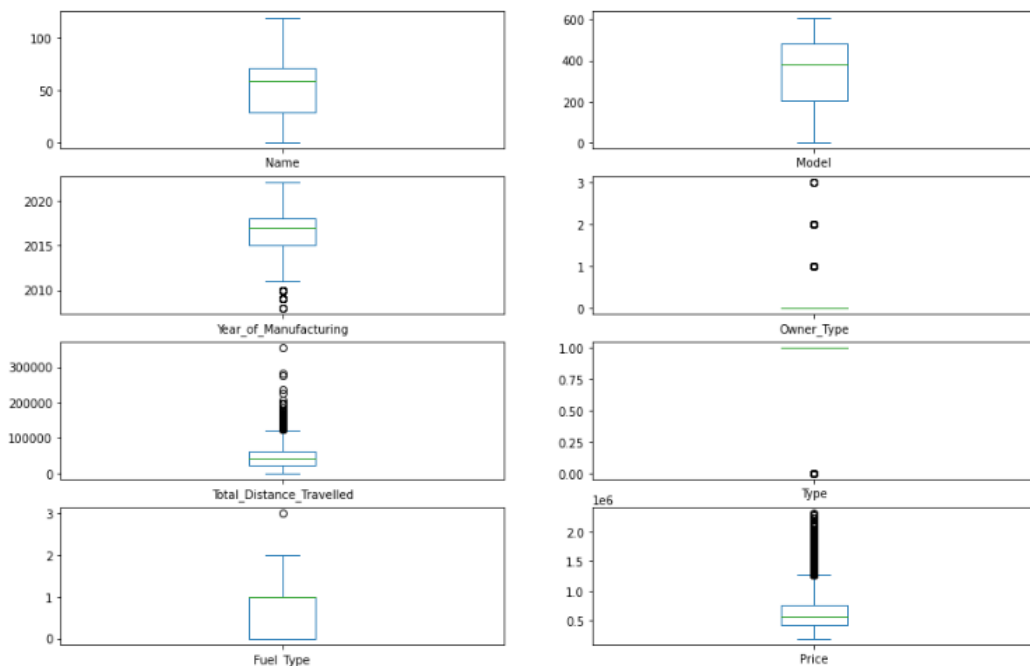
We have encoded our categorical data and find the correlation between each column. As skewness is acceptable so there is no need to scale the skewness.

Removing the Outliers

Plotting the outliers

```
In [61]: df.plot(kind='box',subplots=True,layout=(4,2),figsize=(15,10))
```

```
Out[61]: Name                AxesSubplot(0.125,0.71587;0.352273x0.16413)
         Model              AxesSubplot(0.547727,0.71587;0.352273x0.16413)
         Year_of_Manufacturing AxesSubplot(0.125,0.518913;0.352273x0.16413)
         Owner_Type          AxesSubplot(0.547727,0.518913;0.352273x0.16413)
         Total_Distance_Travelled AxesSubplot(0.125,0.321957;0.352273x0.16413)
         Type                AxesSubplot(0.547727,0.321957;0.352273x0.16413)
         Fuel_Type           AxesSubplot(0.125,0.125;0.352273x0.16413)
         Price               AxesSubplot(0.547727,0.125;0.352273x0.16413)
         dtype: object
```



Removing outliers

```
In [9]: from scipy.stats import zscore
z=np.abs(zscore(df))
threshold=3
print(np.where(z>3))
df_new=df[(z<3).all(axis=1)]
df=df_new
df.shape
```

```
(array([ 3, 9, 12, 14, 23, 29, 32, 34, 65, 67, 74,
       76, 80, 98, 108, 122, 124, 137, 142, 150, 167, 173,
       188, 193, 294, 304, 332, 338, 357, 368, 368, 372, 375,
       376, 397, 413, 417, 424, 437, 452, 514, 538, 552, 555,
       635, 675, 684, 701, 704, 712, 713, 723, 729, 735, 767,
       775, 777, 796, 798, 806, 809, 815, 839, 850, 855, 859,
       901, 964, 974, 989, 1035, 1079, 1081, 1087, 1094, 1098, 1121,
       1122, 1128, 1128, 1136, 1137, 1154, 1199, 1211, 1230, 1297, 1326,
       1363, 1365, 1384, 1388, 1389, 1390, 1398, 1407, 1410, 1417, 1435,
       1442, 1465, 1485, 1493, 1499, 1501, 1515, 1517, 1530, 1553, 1581,
       1638, 1654, 1678, 1679, 1694, 1696, 1699, 1699, 1705, 1708, 1710,
       1736, 1743, 1767, 1778, 1844, 1859, 1863, 1866, 1878, 1970, 2044,
       2046, 2082, 2100, 2107, 2146, 2155, 2182, 2184, 2236, 2247, 2249,
       2255, 2257, 2262, 2285, 2288, 2298, 2311, 2328, 2330, 2338, 2355,
       2364, 2368, 2384, 2386, 2405, 2457, 2463, 2498, 2505, 2519, 2526,
       2527, 2552, 2592, 2609, 2620, 2654, 2662, 2690, 2692, 2693, 2698,
       2700, 2714, 2727, 2749, 2757, 2774, 2810, 2810, 2812, 2836, 2854,
       2876, 2879, 2885, 2887, 2908, 2933, 2938, 2940, 2973, 2993, 2998,
       2999, 3000, 3005, 3010, 3034, 3044, 3052, 3098, 3128, 3169, 3175,
       3177, 3194, 3197, 3199, 3224, 3265, 3267, 3289, 3301, 3305, 3332,
       3337, 3346, 3347, 3352, 3361, 3361, 3371, 3380, 3383, 3388, 3404,
       3405, 3410, 3412, 3420, 3446, 3449, 3453, 3454, 3456, 3457, 3459,
       3459, 3471, 3477, 3481, 3491, 3509, 3513, 3518, 3536, 3539, 3567,
       3600, 3618, 3626, 3636, 3637, 3644, 3648, 3649, 3675, 3698, 3698,
       3700, 3737, 3759, 3785, 3797, 3800, 3806, 3856, 3859, 3870, 3890,
       3909, 3918, 3929, 3941, 3942, 3988, 3991, 4004, 4035, 4056, 4060,
       4061, 4077, 4079, 4096, 4099, 4109, 4125, 4131, 4134, 4136, 4140,
       4185, 4227, 4239, 4308, 4316, 4398, 4433, 4443, 4511, 4548, 4594,
       4640, 4669, 4692, 4698, 4714, 4721, 4725, 4752, 4771, 4772, 4794,
       4906], dtype=int64), array([7, 2, 7, 7, 2, 7, 2, 7, 3, 7, 7, 7, 7, 3,
       3, 7, 7, 3, 7, 7, 7, 4, 3,
       7, 7, 3, 4, 3, 4, 4, 3, 4, 4, 2, 7, 7, 4, 4, 4, 7, 7, 4, 3, 3, 4,
       4, 3, 7, 7, 7, 7, 2, 7, 7, 3, 7, 3, 3, 7, 3, 3, 7, 4, 7, 7, 7, 7,
       7, 4, 4, 3, 3, 7, 4, 4, 4, 7, 4, 4, 3, 4, 7, 2, 3, 4, 7, 3, 4, 4,
       3, 7, 7, 7, 7, 7, 2, 3, 7, 3, 7, 3, 7, 3, 7, 3, 4, 7, 7, 7,
       4, 4, 4, 3, 3, 3, 4, 7, 4, 4, 2, 4, 7, 7, 3, 4, 4, 3, 4, 3,
       7, 7, 3, 7, 4, 7, 7, 4, 7, 7, 7, 3, 4, 7, 7, 4, 3,
       4, 3, 7, 7, 7, 4, 7, 7, 7, 7, 3, 3, 7, 7, 7, 7, 7, 7, 7, 7,
       7, 4, 3, 7, 7, 7, 4, 7, 3, 7, 4, 7, 7, 7, 4, 7, 7, 7, 3, 3,
       7, 3, 4, 3, 3, 7, 4, 4, 7, 7, 7, 7, 7, 7, 3, 7, 3, 2, 3,
       3, 3, 7, 3, 2, 4, 3, 3, 7, 7, 4, 3, 4, 3, 2, 2, 3, 3, 7, 2,
       3, 3, 7, 4, 3, 3, 3, 3, 4, 7, 3, 3, 3, 4, 3, 4, 3, 7, 3, 2, 2, 6,
       7, 7, 7, 7, 3, 7, 4, 4, 7, 7, 7, 3, 4, 7, 3, 3, 3, 7, 4, 3, 4,
       7, 7, 4, 4, 3, 3, 3, 3, 4, 3, 3, 3, 3, 7, 7, 7, 4, 4, 3, 3,
       4, 7, 7, 7, 4, 3, 7, 3, 3, 4, 3, 4], dtype=int64))
```

Out[9]: (4670, 8)

We have some outliers present in the dataset, so let's handle them also. As the outliers in the dataset will affect our ML model. We need to remove all the outliers present in the dataset.

There is something called zscore which indicates how many standard deviations away an element is from the mean. We consider the points as outliers whose zscore is above 3 or less than -3. So we need to remove all such points from our dataset.

Using the threshold, we have removed all the points where the zscore is greater than 3. Now the total number of rows after removing the outliers are 4670.

MODEL BUILDING

We will import important libraries for the building the ML model and defining the different models for our easiness.

Finding the best random state for the train test split.

Model Building

```
In [13]: #importing the different machine Learning models

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import r2_score
```

```
In [14]: # defining the different models

lg=LinearRegression()
rdr=RandomForestRegressor()
svr=SVR()
dtr=DecisionTreeRegressor()
knr=KNeighborsRegressor()
```

Finding the best random state

```
In [15]: model=[lg,rdr,svr,dtr,knr]
maxAcc=0
maxRS=0
for i in range(40,60):
    x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=i,test_size=.20)
    lg.fit(x_train,y_train)
    pred=lg.predict(x_test)
    acc=r2_score(y_test,pred)
    if acc>maxAcc:
        maxAcc=acc
        maxRS=i
print('Best Accuracy score is', maxAcc , 'on random state', maxRS)

Best Accuracy score is 0.4459634857166468 on random state 45
```

```
In [16]: x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=45,test_size=.20)
```

Regression Algorithms


We have use five different regression algorithms to find the best model for our problem.

1. Linear Regression

Linear regression is one of the easiest and most popular Machine Learning algorithms. It is a statistical method that is used for predictive analysis. Linear regression makes predictions for continuous/real or numeric variables such as **sales, salary, age, product price**, etc.

Linear regression algorithm shows a linear relationship between a dependent (y) and one or more independent (x) variables, hence called as linear regression. Since linear regression shows the linear relationship, which means it finds how the value of the dependent variable is changing according to the value of the independent variable.

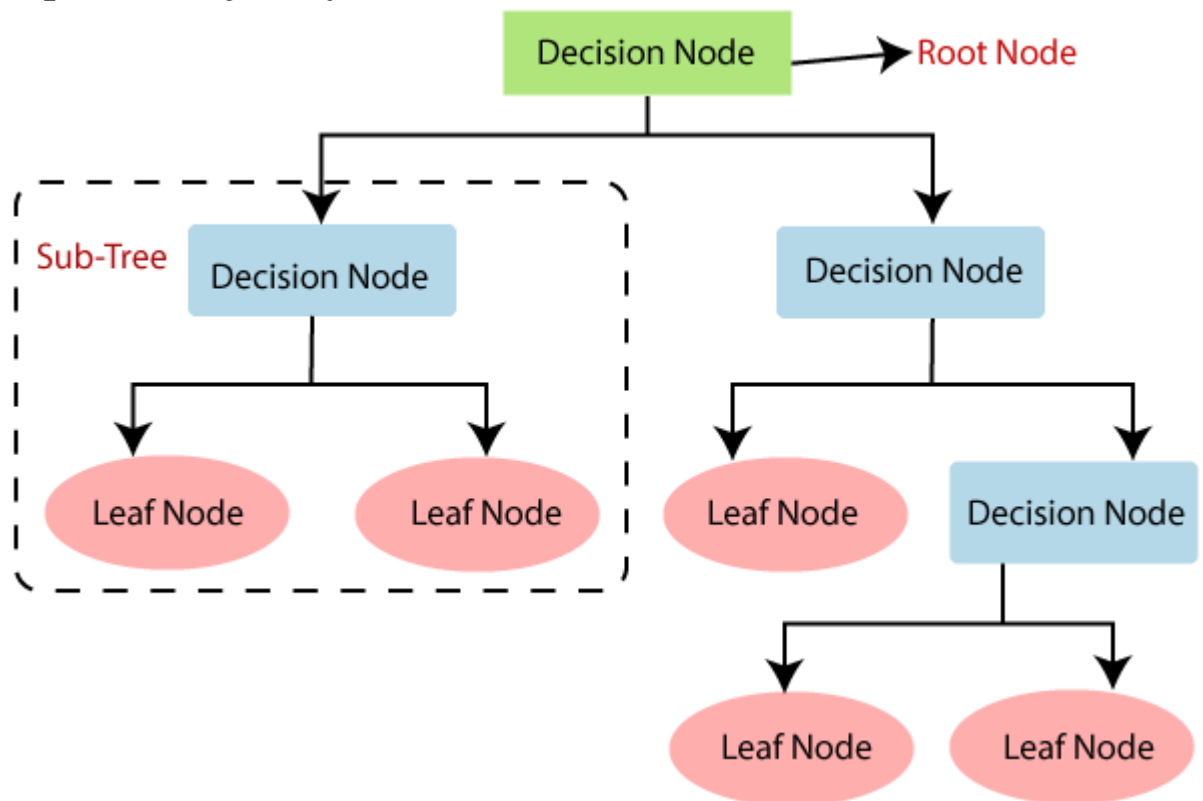
The linear regression model provides a sloped straight line representing the relationship between the variables.



2. Decision Tree Regressor


Decision Tree can be used both in classification and regression problem. Decision trees are predictive models that use a set of binary rules to calculate a target value. Each individual tree is a fairly simple model that has branches, nodes and leaves. Decision tree regression observes features of an object and **trains a model in the structure of a tree to predict data in the future to produce meaningful continuous output**. Continuous output

means that the output/result is not discrete, i.e., it is not represented just by a discrete, known set of numbers or values.



3. Support Vector Regressor


Support Vector Regression is a supervised learning algorithm that is used to predict discrete values. Support Vector Regression uses the same principle as the SVMs. The basic idea behind SVR is to find the best fit line. In SVR, the best fit line is the hyperplane that has the maximum number of points.



4. K Neighbor Regression

KNN regression is a non-parametric method that, in an intuitive manner, approximates the association between independent variables and the continuous outcome by averaging the observations in the same *neighbourhood*. The size of the neighbourhood needs to be set by the analyst or can be chosen using cross-validation (we will see this later) to select the size that minimises the mean-squared error.


kNN Regression



5. Random Forest Regression

A random forest is a machine learning technique that's used to solve regression and classification problems. It utilizes ensemble learning, which is a technique that combines many classifiers to provide solutions to complex problems. Random Forest works in two-phase first is to create the random forest by combining N decision tree, and second is to make predictions for each tree created in the first phase.

Random Forest Simplified



Libraries Used for Regression Models

- Linear Regression
 - `from sklearn.linear_model import LinearRegression`
- Decision Tree Regressor
 - `from sklearn.tree import DecisionTreeRegressor`
- Support Vector Regressor
 - `from sklearn.svm import SVR`
- Kneighbor Regressor
 - `from sklearn.neighbors import KNeighborsRegressor`
- Random Forest Regressor
 - `from sklearn.ensemble import RandomForestRegressor`

Model Accuracy:

| MODEL | ACCURACY |
|--------------------------|----------------------|
| Linear Regression | 0.4459634857166468 |
| Decision Tree Regressor | 0.8856552523096618 |
| Support Vector Regressor | -0.08005490482792754 |
| Kneighbor Regressor | 0.04292499937121386 |
| Random Forest Regressor | 0.8815066126350098 |

Linear Regression

```
In [17]: lg.fit(x_train,y_train)
pred1=lg.predict(x_test)
acc=r2_score(y_test,pred1)
print('Accuracy Score: ',acc)

Accuracy Score: 0.4459634857166468
```

Decision Tree Regressor

```
In [18]: dtr.fit(x_train,y_train)
pred2=dtr.predict(x_test)
acc=r2_score(y_test,pred2)
print('Accuracy Score: ',acc)

Accuracy Score: 0.8856552523096618
```

Support Vector Regressor

```
In [19]: svr.fit(x_train,y_train)
pred3=svr.predict(x_test)
acc=r2_score(y_test,pred3)
print('Accuracy Score: ',acc)

Accuracy Score: -0.08005490482792754
```

K Neighbor Regressor

```
In [20]: knr.fit(x_train,y_train)
pred4=knr.predict(x_test)
acc=r2_score(y_test,pred4)
print('Accuracy Score: ',acc)

Accuracy Score: 0.04292499937121386
```

Random Forest Regressor

```
In [21]: rdr.fit(x_train,y_train)
pred5=rdr.predict(x_test)
acc=r2_score(y_test,pred5)
print('Accuracy Score: ',acc)

Accuracy Score: 0.942351061631132
```

Hence, we are getting the best accuracy score through the Random Forest Classifier Model. We will go ahead with this to find the cross val score and hypermeter tuning.

Cross Val Score & Hypermeter Tuning

Cross-validation provides information about how well a classifier generalizes, specifically the range of expected errors of the classifier. Cross Val Score tells how the model is generalized at a particular cross validation.

At CV=6 we get the best results i.e. the Random Forest Classifier more generalized at cv=6, so we calculate the hyper parameters at this value.

We will find which parameters of random forest classifier are the best foe our model. We will do this using Grid Search CV method & also calculate the accuracy score at those best parameters.

Cross Val Score

```
In [22]: from sklearn.model_selection import cross_val_score
for i in range(3,7):
    cr=cross_val_score(rndr,x,y,cv=i)
    cr_mean=cr.mean()
    print("at cv= ", i)
    print('cross val score = ',cr_mean*100)

at cv= 3
cross val score = 85.942859785015
at cv= 4
cross val score = 90.8769072158919
at cv= 5
cross val score = 91.41279286544679
at cv= 6
cross val score = 91.7821441798227
```

Hypermeter Tuning

```
In [23]: from sklearn.model_selection import GridSearchCV
# creating parameters
para={'criterion':['squared_error','absolute_error','poisson'],
      'max_features':['sqrt','log2'],
      'max_depth':[1,2,3,4,5]}

GCV=GridSearchCV(rndr,para,cv=6,scoring='accuracy')
GCV.fit(x_train,y_train)
GCV.best_params_
```

```
Out[23]: {'criterion': 'squared_error', 'max_depth': 1, 'max_features': 'sqrt'}
```

```
In [24]: GCV_pred=GCV.best_estimator_.predict(x_test)
r2_score(y_test,GCV_pred)
```

```
Out[24]: 0.184987654844048
```

For Random Forest we are not getting very good results. The r^2 score is very low for the best parameters. So we will move to the next best model and find the best parameters for Decision Tree Regressor and calculate the accuracy.

Cross val Score

```
In [26]: from sklearn.model_selection import cross_val_score
for i in range(3,7):
    cr=cross_val_score(dtr,x,y,cv=i)
    cr_mean=cr.mean()
    print("at cv= ", i)
    print('cross val score = ',cr_mean*100)

at cv= 3
cross val score = 76.33747044735803
at cv= 4
cross val score = 85.46379663489655
at cv= 5
cross val score = 87.27113441425047
at cv= 6
cross val score = 87.28937269229745
```

Hypermeter Tuning

```
In [28]: from sklearn.model_selection import GridSearchCV
# creating parameters
para={'criterion':['squared_error','absolute_error','poisson','friedman_mse'],
      'max_features':['sqrt','log2','auto'],
      'max_depth':[1,2,3,4,5],
      'splitter':['best','random']}

GCV=GridSearchCV(dtr,para,cv=6,scoring='accuracy')
GCV.fit(x_train,y_train)
GCV.best_params_

Out[28]: {'criterion': 'absolute_error',
          'max_depth': 5,
          'max_features': 'auto',
          'splitter': 'random'}

In [29]: GCV_pred=GCV.best_estimator_.predict(x_test)
r2_score(y_test,GCV_pred)

Out[29]: 0.4663703959145542
```

With decision tree regressor we are getting a better accuracy score so we will save the decision tree model

We are getting the better accuracy for the decision tree Regressor, so we will save this model for our future predictions.

Saving the Model

Saving the best model – Decision Tree Regression in this case for future predictions. Let's see what are the actual test data and what our model predicts.

Saving the Model

```
In [32]: import pickle
filename='used_car_price.pkl'
pickle.dump(dtr, open(filename, 'wb'))
```

Conclusion

```
In [31]: a=np.array(y_test)
pred=np.array(pred2)
Used_Car_Price=pd.DataFrame({'Actual':a,'Predicted':pred})
Used_Car_Price
```

| | | |
|----|---------|-----------|
| 3 | 607499 | 476349.0 |
| 4 | 557399 | 565499.0 |
| 5 | 770199 | 720999.0 |
| 6 | 421399 | 428999.0 |
| 7 | 1097099 | 1097099.0 |
| 8 | 340499 | 340499.0 |
| 9 | 490199 | 490199.0 |
| 10 | 436699 | 415099.0 |
| 11 | 863599 | 863599.0 |
| 12 | 723699 | 846699.0 |
| 13 | 407999 | 292599.0 |
| 14 | 522999 | 460599.0 |
| 15 | 818599 | 1055199.0 |

Hence up to some good extensions our model predicted so well.

CONCLUSION

Conclusion of the Study

The results of this study suggest following outputs which might be useful for the client to evaluate the price on the basis of new data:

- There are lot of things that is going to decide the price of a car. As we see above in our visualizations, a lot of things affect the price like year of manufacturing, brand, distance travelled & many more. One needs to analyse every aspect to have good hands on the prediction of the price.
 - With the machine learning it become easier to predict the price but yes it is not 100% accurate, it provides an idea and accordingly we can analyse the market and prepare the strategies to grab the opportunities.
-
- **Learning Outcomes of the Study in respect of Data Science**
 - As our first step to collect the fresh or new data for the project in accordance with today's trend, scraping of data is not an easy task but yes every difficulty teaches something & I learn so much exceptions handling and error handling during the first phase of project.
 - Second phase is to create the machine learning model for the client to have a good hands on the prediction after the Covid 19 situation, which is an interesting part. Data pre-processing, null values handling, data cleaning and data engineering every time handle in a different way which is a good learning part.
 - Every model whose accuracy is higher doesn't mean it predicts always best. As we see we are getting the Random Forest accuracy highest but its hyperparameters accuracy is very low so we have to double check on with another algorithm as well.